

Predictive Modeling and Feature Selection: An Analysis of the Ames Housing Dataset

Subhayan Biswas

Roll no.: 712

MDTS 4214 - Final Assignment

10/04/2026

St. Xavier's College (Autonomous), Kolkata

Introduction & Data Preprocessing

Objective: Build optimized prognostic (pricing) and diagnostic (classification) models using the Ames Housing dataset.


Data Cleaning & Stability Metrics:

- **Initial State:** 2930 observations, 82 explanatory variables (39 are Quantitative and 43 are Qualitative).
- **Dummy Encoding:** Handled categorical variables, expanding the matrix to 257 features.
- **Multicollinearity Resolution:** Dropped perfectly collinear variables (e.g., Total Bsmt SF, Garage Qual_None) to resolve $VIF = \infty$.
- **Outlier Removal:** Removed 195 influential points (Cook's D $> 4/n$) and 44 extreme outliers ($|Std. Resid| > 3$).

Univariate Analysis: Top Predictors

Methodology: Fitted 257 Simple Linear Regression models to isolate the individual predictive power of each feature.

Top 5 Univariate Predictors vs. SalePrice



Univariate Analysis: Interpretations


Top 5 Individual Predictors (R^2):

1. **Overall Qual:** 0.6388
2. **Gr Liv Area:** 0.4995
3. **Garage Cars:** 0.4193
4. **Total Bsmt SF:** 0.4001
5. **Exter Qual_TA:** 0.3492

Interpretation: Quality and scale (Square Footage) drive the market. Overall Qual alone explains almost 64% of price variance, validating its heavy weighting in subsequent multivariate models.

Feature Selection: Forward Stepwise


Goal: Compare subset selection strategies to bypass the curse of dimensionality.



Result: Stepwise selection greedily drafted 85 columns. Despite a high Adjusted R^2 (0.9406), the AIC (61,217) indicates severe overfitting and unnecessary model complexity.

Feature Selection: LASSO Shrinkage



Goal: Apply an L_1 penalty to shrink redundant noise and find the optimal subset.



Result: LASSO ($\alpha = 1200$) isolated exactly **12 features**. The model achieved an AIC of 50,810. The lower AIC proves LASSO achieved a vastly superior Bias-Variance tradeoff compared to Stepwise selection.

Regularization Validation: Ridge & Elastic Net


Validation: We applied L_2 and Hybrid penalties to the final 12 LASSO features to stress-test for residual multicollinearity.



Interpretation: Ridge ($R^2 = 0.9119$) and Elastic Net ($R^2 = 0.9118$) performed nearly identically to unpenalized OLS. Elastic Net automatically selected an L_1 ratio of 1.0, mathematically proving that pure LASSO was the optimal feature selector.

Prognostic Modeling: OLS vs. KNN


Methodology: Predicting continuous SalePrice using the 12 optimal features.



Interpretation: Multiple Linear Regression ($R^2 = 0.9119$) slightly outperformed KNN Regression ($R^2 = 0.8885$). Housing pricing metrics scale highly linearly, making standard OLS coefficient weighting highly effective and robust.


Diagnostic Modeling: Binary Classification

Methodology: Engineered a binary target ("Standard" vs. "Expensive" based on median price) using Logistic Regression on the 12 selected features.



Diagnostic Modeling: Multi-Category Classification

Methodology: Categorized the market into three tiers (Budget, Standard, Luxury).



Final Project Takeaways:

- **Dimensionality Reduction:** Shrinkage methods (LASSO) successfully reduced 257 dimensions to 12 ($\sim 5\%$), entirely eliminating the noise captured by Forward Stepwise selection.
- **Predictive Power:** The sparse 12-feature matrix explained $\sim 91\%$ of the continuous pricing variance ($\text{RMSE} \approx 22,600$).
- **Model Versatility:** The selected features proved mathematically robust across both prognostic algorithms (OLS/KNN) and diagnostic classification algorithms (Logistic), confirming the underlying stability of the Ames real estate market signals.